

DISTRIBUTED LLM INFERENCE EFFICIENCY: THROUGHPUT VS. ENERGY DISSIPATION

Microarchitectural Benchmark on 8x H100 GPU Cluster (70B Model, FP8 Tensor Cores)




Figure 1: Measured execution metrics across varying batch workloads. Shaded bands depict ± 1 sigma experimental variance over 50 iterations.